



CS-671 DEEP LEARNING AND APPLICATIONS

ASSIGNMENT 5

Convolutional Neural Networks

Submitted by:

Group 13

Prachi Sharma (V21078)

Ayush Tiwari (T22053)

Instructor:

Prof. A D Dileep

May 1, 2023

Table of Contents

Convolutional Neural Networks.....	2
1. The convolution operation.....	3
2. Structure of a CNN.....	3
3. Backpropagation in CNNs.....	3
Experiments, Results.....	4
and Analysis.....	4
1. Classification using different CNNs.....	5
a. Architecture-1 (4 layer architecture).....	6
b. Architecture-2 (5 layer architecture).....	7
c. Architecture-3 (6 layer architecture).....	8
d. Results of best architecture.....	9
2. Classification using VGG.....	13

Convolutional Neural Networks

1. The convolution operation

The convolutional neural networks, abbreviated as CNNs, are a special type of neural network used to perform learning on data that has a grid-like structure. A CNN uses a *convolution* operation in at least one of its layers. A convolution operation on a 1-D input is defined as:

$$s_t = \sum_{a=0}^{\infty} x_{t-a} w_a = (\mathbf{x} * \mathbf{w})_t$$

Here, \mathbf{x} is the **input**, \mathbf{w} is the **kernel** or filter and t is the time step. The output \mathbf{s} is called the feature map.

2. Structure of a CNN

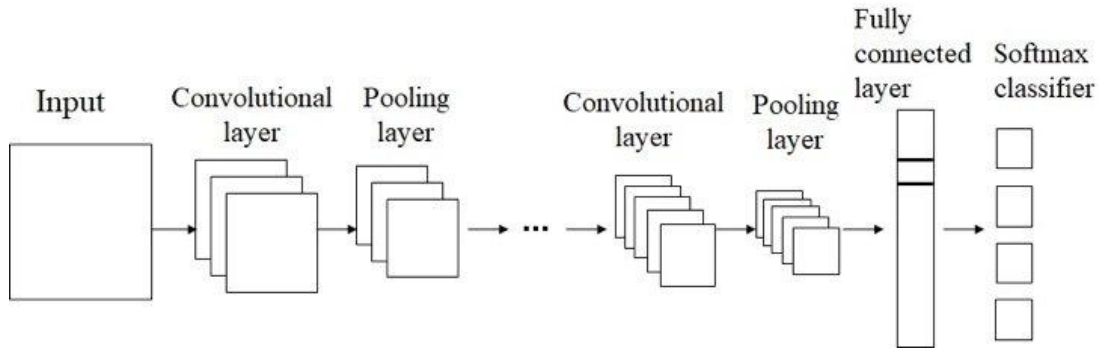


Fig 1: Structure of a basic CNN

A typical CNN has a structure as shown in Fig 1. The input is sent to a **convolutional layer**, which has a K number of filters of size F . A convolution operation is performed by all the filters on the input and the output is then sent to a pooling layer. A **pooling layer** modifies the outputs of the previous layer by operating it with a pooling function. Two such functions are *maxpool* and *average pool*. A max pool function reports the maximum output within a rectangular output while an average pool takes the average of all values in a rectangular output. After several consecutive convolution and pooling layers, the network is connected to a dense layer of neuron units, called **dense layers**, and finally a fully-connected **output layer**.

3. Backpropagation in CNNs

The connections in a convolution neural network are much sparser and weights are shared to make learning easier. CNNs can be trained just like a feed-forward neural network by backpropagating the error. Only the active parameter weights are updated at each instance of backpropagation.

Experiments, Results and Analysis

In this section, we will present our results and observations on image classification using different CNN architectures. In the first part, we train three different CNNs to classify a subset of the Caltech-101 dataset, which consists of color images. Different performance measures have been observed and compared of all three networks to conclude the best performing network. In the second part, we make use of the pretrained VGG19 model available in Tensorflow and perform transfer learning to classify the given image dataset.

1. Classification using different CNNs

We have been given 5 classes from 101 classes of the Caltech-101 image dataset. Each class has 80 images belonging to classes: **brain, buddha, ketch, Leopards, scorpion**. The train, validation and test split among all classes is 5:1:2. To ensure same input dimension, the images have been resized to 224×224 .

Following architectures have been considered for the assigned task.

Model	Input layer	Convolution layer 1	Pooling layer 1	Convolution layer 2	Pooling layer 2	Convolution layer 3	Pooling layer 3	Convolution layer 4	Pooling layer 4	Flatten layer	Dense layer	Output layer
1	$28 \times 28 \times 3$	$F = 11$ $K = 8$ $S = 4$ $P = 0$	$F = 3$ $K = 8$ $S = 2$ $P = 0$	$F = 5$ $K = 16$ $S = 1$ $P = 0$	$F = 3$ $K = 16$ $S = 2$ $P = 0$	-	-	-	-		128	5
2	$28 \times 28 \times 3$	$F = 11$ $K = 8$ $S = 4$ $P = 0$	$F = 3$ $K = 8$ $S = 2$ $P = 0$	$F = 5$ $K = 16$ $S = 1$ $P = 0$	$F = 3$ $K = 16$ $S = 2$ $P = 0$	$F = 3$ $K = 32$ $S = 1$ $P = 0$	$F = 3$ $K = 32$ $S = 2$ $P = 0$	-	-		128	5
3	$28 \times 28 \times 3$	$F = 11$ $K = 8$ $S = 4$ $P = 0$	$F = 3$ $K = 8$ $S = 2$ $P = 0$	$F = 5$ $K = 16$ $S = 1$ $P = 0$	$F = 3$ $K = 16$ $S = 2$ $P = 0$	$F = 3$ $K = 32$ $S = 1$ $P = 0$	-	$F = 3$ $K = 64$ $S = 1$ $P = 0$	$F = 3$ $K = 64$ $S = 2$ $P = 0$		128	5

The output layer has 5 hidden units representing five classes of given dataset. The data is normalized using MinMax normalization to scale all input features in the range of $[0,1]$.

Parameters for all models:

Parameters	Values
Hidden layers	Relu activation
Output layer	Softmax activation
Loss function	Cross Entropy
Batch size	32
Learning rate (η)	0.001
Optimiser	Adam
Beta 1 (β_1)	0.9
Beta 2 (β_2)	0.999
Epsilon (ϵ)	10^{-8}
Convergence criterion	Difference between successive epochs is less than or equal to 10^{-4} .

a. Architecture-1 (4 layer architecture)

Layer Name	Output dimensions
Input layer	(224,224,3)
Convolution layer 1	(54,54,8)
Max Pooling layer 1	(26,26,8)
Convolution layer 2	(22,22,16)
Max Pooling layer 2	(10,10,16)
Flatten	1600
Dense layer 1	128
Output layer	5

Confusion matrix on validation data	ACTUAL					
	P R E D I C T E D	9	0	0	0	1
		0	8	1	1	0
		0	2	8	0	0
		0	0	0	10	0
		0	0	0	0	10
Classification accuracy	Train			Validation		
	100%			90%		

b. Architecture-2 (5 layer architecture)

Layer Name	Output dimensions
Input layer	(224,224,3)
Convolution layer 1	(54,54,8)
Max Pooling layer 1	(26,26,8)
Convolution layer 2	(22,22,16)
Max Pooling layer 2	(10,10,16)
Convolution layer 3	(8,8,32)
Max Pooling layer 3	(3,3,32)
Flatten	288
Dense layer 1	128
Output layer	5

Confusion matrix on validation data	A C T U A L				
	P R E D I C T E D	10	0	0	0
		0	8	0	2
		1	1	7	0
		0	0	0	10
		1	0	0	1
Classification accuracy	Train			Validation	
	100%			86%	

c. Architecture-3 (6 layer architecture)

Layer Name	Output dimensions
Input layer	(224,224,3)
Convolution layer 1	(54,54,8)
Max Pooling layer 1	(26,26,8)
Convolution layer 2	(22,22,16)
Max Pooling layer 2	(10,10,16)
Convolution layer 3	(8,8,32)
Convolution layer 4	(6,6,64)
Max Pooling layer 4	(2,2,64)
Flatten	256
Dense layer 1	128
Output layer	5

Confusion matrix on validation data	ACTUAL					
	P R E D I C T E D	8	1	0	0	1
		1	7	0	2	0
		1	1	8	0	0
		0	0	1	9	0
		0	0	0	1	9
Classification accuracy	Train			Validation		
	100%			82%		

d. Results of best architecture

After comparing the classification accuracy on validation data, we can observe that architecture 1 with 4 hidden layers works best. This section provides results on best architecture and feature map visualization for the same.

- Results on test data

Confusion matrix	ACTUAL					
	P R E D I C T E D	18	2	0	0	0
		1	17	1	1	0
		1	3	16	0	0
		0	1	0	19	0
		2	0	1	1	16
Classification accuracy	86%					

- **Feature maps for first convolution layer**




Fig 2: Image of a Scorpion from Caltech-101 dataset




Fig 3: Feature maps for 8 filters of size 11×11 applied on the above image of scorpion.

- **Feature maps for second convolution layer**


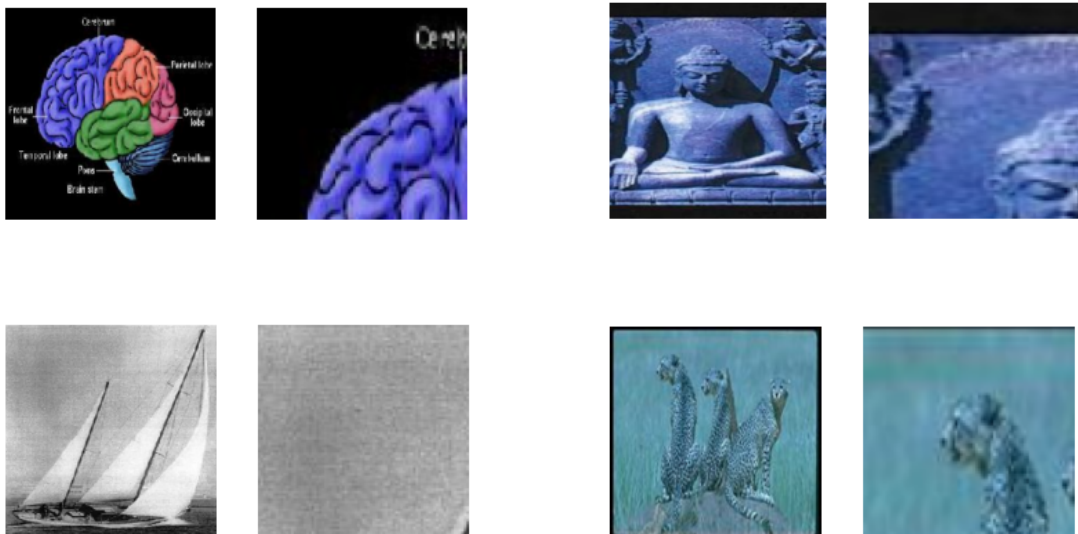


Fig 3: Selected 8 feature maps out of 16 filters of size 5×5 applied on the above feature maps.

- **Visualizing input image patch for maximally activated neuron**

To observe which part of the input image is responsible for maximal activation of our last convolution layer, we backtrack from maximally activated neuron through all previous layers.





Observations:

1. In all the architectures, the classification accuracy on unseen data is not very high ($\leq 90\%$) while the accuracy on training data for is 100% which indicates that the models are not able to generalize well. Lack of training data could be one reason behind it.
2. The feature maps plot of both convolution layers shows how the model is trying to learn different features at different layers. First layer learns simpler features like edges and gradients. Some parts of the feature maps are more activated for a particular color (7th feature map is most activated for blue color). Second layer learns more abstract features, mostly curves and parts of the image.

2. Classification using VGG

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv4 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv4 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv4 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten_1 (Flatten)	(None, 25088)	0
dense_1 (Dense)	(None, 5)	12544
Total params: 20,149,829		
Trainable params: 125,445		
Non-trainable params: 20,024,384		

- Results on train data

	ACTUAL
--	--------

Confusion matrix	P R E D I C T E D	50	0	0	0	0
		0	50	0	0	0
		0	0	50	0	0
		0	0	0	50	0
		0	0	0	0	50
Classification accuracy	100%					

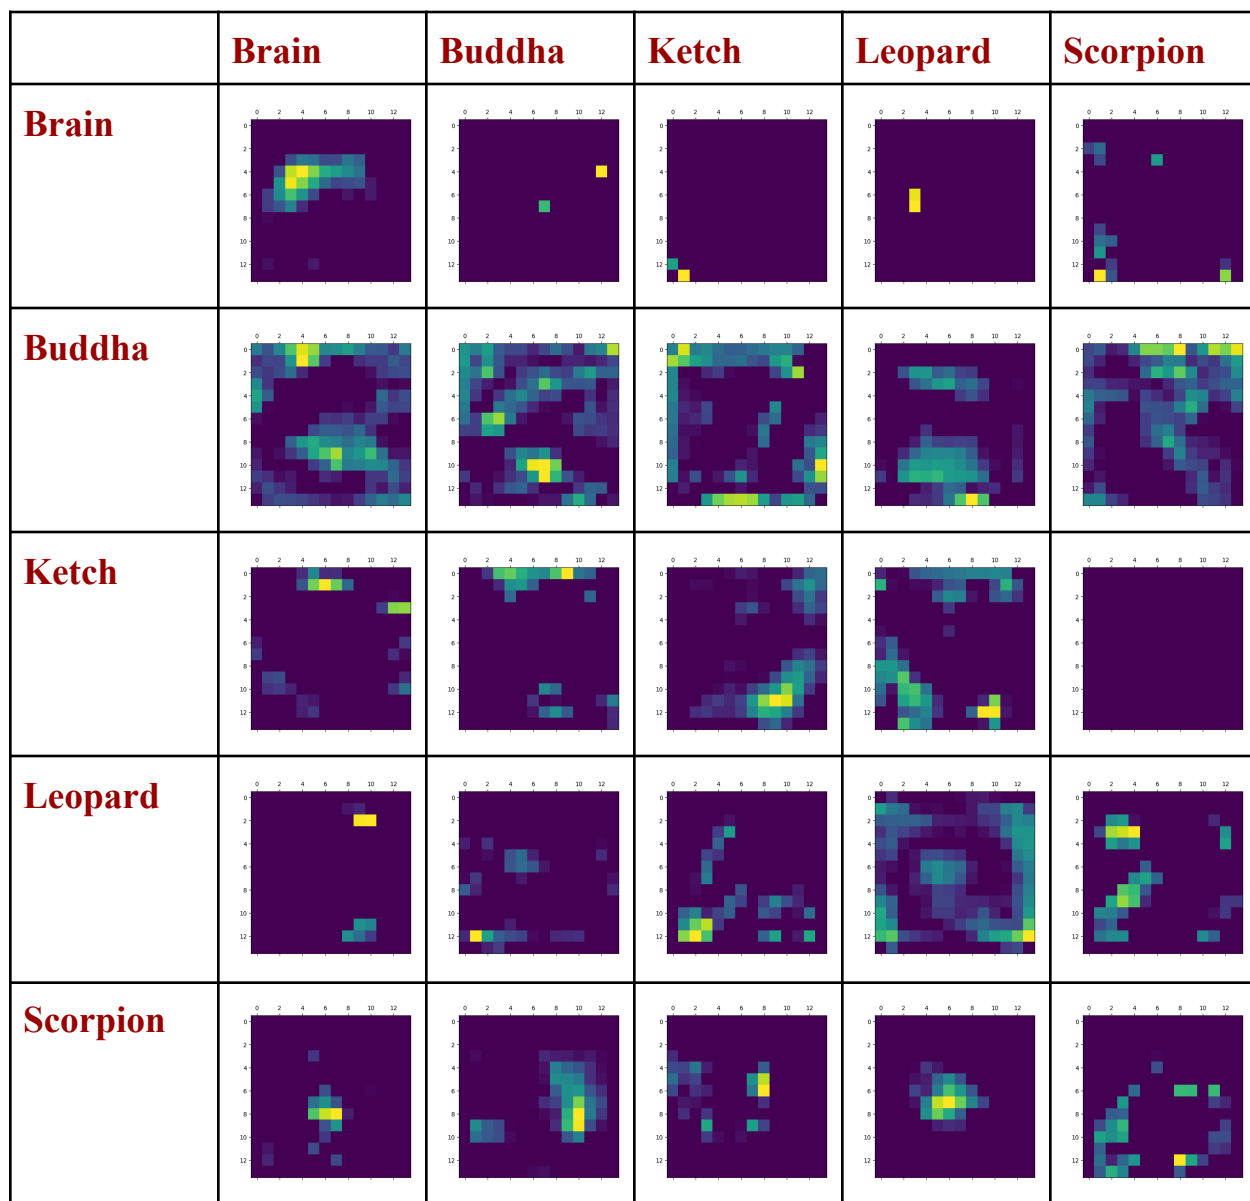
- Results on validation data

Confusion matrix	A C T U A L					
	P R E D I C T E D	10	0	0	0	0
		0	10	0	0	0
		0	0	10	0	0
		0	0	0	10	0
		0	0	0	0	10
Classification accuracy	100%					

- Results on test data

Confusion matrix	A C T U A L					
	P R E D I C T E D	20	0	0	0	0
		0	20	0	0	0
		0	0	20	0	0
		0	0	0	20	0
		0	0	0	0	20
Classification accuracy	100%					

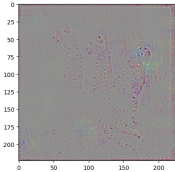
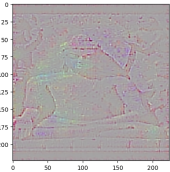
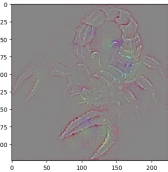
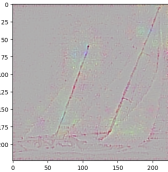
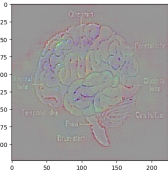
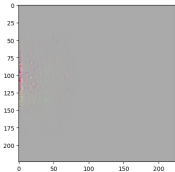
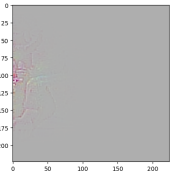
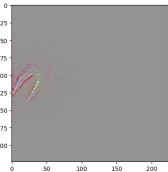
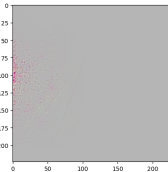
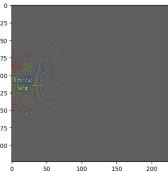
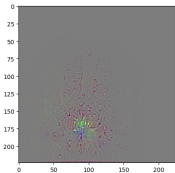
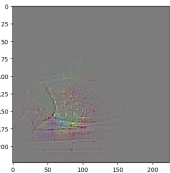
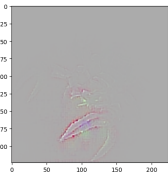
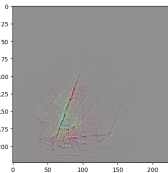
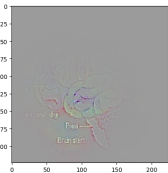
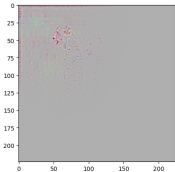
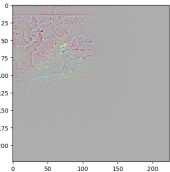
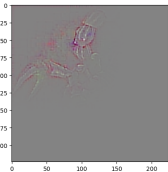
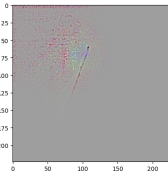
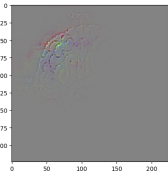
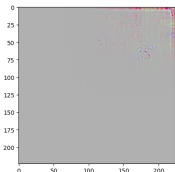
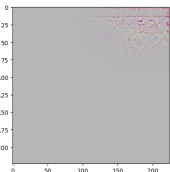
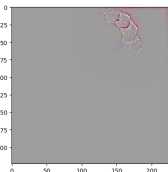
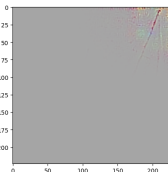
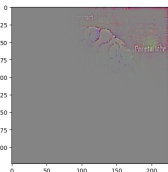
GradCam Results:



Observation:

If we consider the diagonal heatmaps of the above classes, then it somewhat represents the original shape which the model is trying to learn. The grid (1,1) clearly represents the shape of the brain. The other non-diagonal value is what the heatmap corresponds to the features which the model learns to distinguish other classes from that particular class.

Using Guided BackProp:

Neuron number	Leopard	Buddha	Scorpion	Ketch	Brain
93619					
43517					
67482					
23476					
5987					

Observations:

The first neuron (93619) is the maximally activated neuron among all other neurons. On tracing back to the filter, it can be clearly seen that this particular neuron (and its associated feature maps in the previous layers) is responsible for

learning a better overall representation of each of the classes. The other neurons are randomly selected and it is observed that each neuron tries to learn some of the patch/portion of the input images.